

HW2

Problem 1: Apply the following models on the Fashion Mnist Dataset. Train the model with the training data and evaluate the model with the test data.

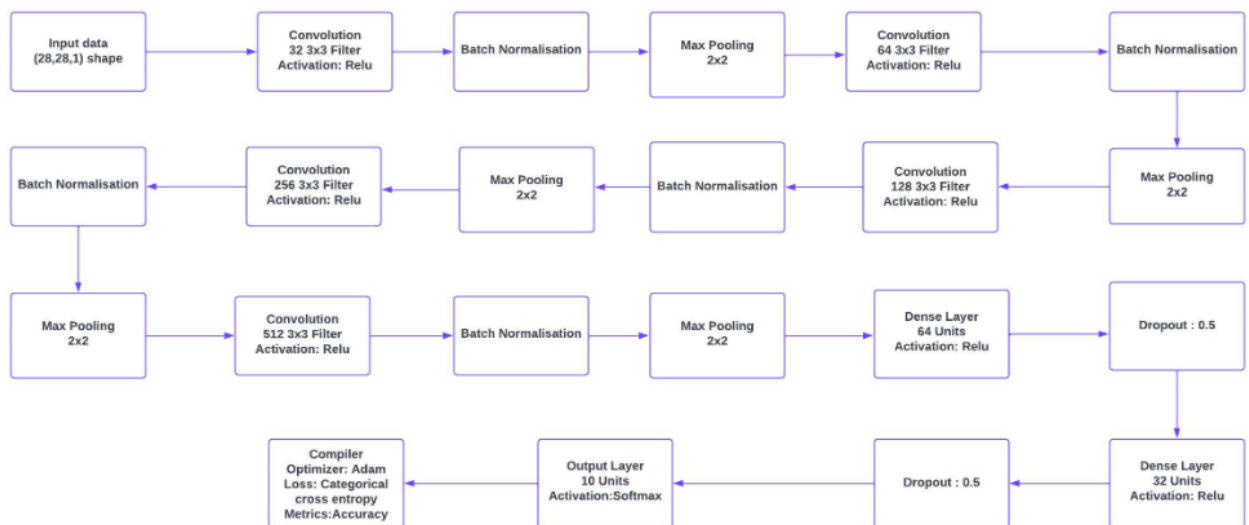
a: CNN model from scratch: Develop a CNN model with 5 convolutional layers (with kernel size= 3, stride = 1, padding = “same”, activation function = “relu” (with following Max Pooling layer (Size= 2) and 3 fully connected layer (including one output layer). After each of the Convolutional layer apply Batch Normalization. In the fully connected layer apply dropout.

Ans:

Model description, summary and hyper-parameters used are given below.

For detailed description of the model, accuracy and prediction results, please refer to the notebook.

Model Architecture



1. The input shape is (28,28,1), which means that the input image has a height and width of 28 pixels, and a depth of 1 channel (grayscale).
2. This model is a convolutional neural network (CNN) architecture. It consists of 5 convolutional layers, each followed by a batch normalization layer and a max pooling layer.
3. The output from the convolutional layers is flattened and then fed into two fully connected layers, each with a ReLU activation function and dropout regularization.
4. The final layer is a dense layer with a softmax activation function to produce the output probabilities.
5. The model uses the Adam optimizer and categorical cross-entropy loss function to train, and accuracy is used as the evaluation metric.
6. Model is fitted on the training data with 20 epochs and batch size 32 and validation and training data accuracy and loss is compared to see that model is not overfitting

Hyper-parameters Used

Number of Convolution Layers: 5						
Convolutions Layers	Filter Size	Stride	Number of Filters	Padding	Activation Function	
	Convo Layer 1	3x3	1	32	Same	Relu
	Convo Layer 2	3x3	1	64	Same	Relu
	Convo Layer 3	3x3	1	128	Same	Relu
	Convo Layer 4	3x3	1	256	Same	Relu
	Convo Layer 5	3x3	1	512	Same	Relu
Number of Hidden Layers: 2						
Hidden Layers	Neurons			Activation Function		
Hidden Layer 1	64			Relu		
Hidden Layer 2	32			Relu		
Output Layer						
Output Layer	Neurons			Activation Function		
	10			Softmax		
Regularisation						
Part a		Dropout p = 0.5				
Part b		Dropout p = 0.5 , Data Augmentation				
Model Compilation						
Loss	Optimizer	Learning Rate	Evaluation Metric			
Categorical Cross entropy	Adam	0.001	Accuracy			
Training						
Epochs	Batch Size			Validation Split		
20	32			0.2		
Weight Initialiser		Bias Initialiser				
Be default Glorot		By default zeros				

Model Summary

```
1 modela.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 32)	320
batch_normalization (Batch Normalization)	(None, 28, 28, 32)	128
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 14, 14, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_2 (Conv2D)	(None, 7, 7, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 7, 7, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
conv2d_3 (Conv2D)	(None, 4, 4, 256)	295168

batch_normalization_3 (Batch Normalization)	(None, 4, 4, 256)	1024
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 256)	0
conv2d_4 (Conv2D)	(None, 2, 2, 512)	1180160
batch_normalization_4 (Batch Normalization)	(None, 2, 2, 512)	2048
max_pooling2d_4 (MaxPooling2D)	(None, 1, 1, 512)	0
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 64)	32832
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 32)	2080
dropout_1 (Dropout)	(None, 32)	0
dense_2 (Dense)	(None, 10)	330

```
=====
Total params: 1,607,210
Trainable params: 1,605,226
Non-trainable params: 1,984
```

The model has a total of 1,607,210 parameters, out of which 1,605,226 are trainable and 1,984 are non-trainable. The model consists of convolutional layers, batch normalization layers, max pooling layers, fully connected layers with dropout, and an output layer with softmax activation. The architecture progressively reduces the spatial dimensions of the input, from 28x28 to 1x1, while increasing the number of channels in each convolutional layer. The model was compiled with the Adam optimizer (it computes adaptive learning rates for each parameter and performs both momentum and RMSprop style updates for faster convergence), accuracy as the evaluation metric. and categorical cross-entropy loss function, and trained with accuracy as the evaluation metric.

In total, this model has 1,607,210 parameters, out of which 1,605,226 are trainable.

```
1 end_timea = time.time()
```

```
1 total_timea = end_timea - start_timea
```

```
1 print("Total execution time: ", total_timea, " seconds")
```

```
Total execution time: 200.49847745895386 seconds
```



Evaluation

```
1 # evaluating the performance model on test data
2 test_evala = modela.evaluate(test_X, test_Y_one_hot, verbose=0)
```


```
1 # printing the loss and accuracy value on test data
2 print('Test loss:', test_evala[0])
3 print('Test accuracy:', test_evala[1])
```

Test loss: 0.4806569218635559

Test accuracy: 0.9092000126838684



Confusion Matrix



```
1 print("Precision Score is ", precision_score(test_Y, predictiona, average = 'weighted'))
```

Precision Score is 0.9108256205924419

```
1 print("Recall Score is ", recall_score(test_Y, predictiona, average = 'weighted'))
```

Recall Score is 0.9077

```
1 print("F1 Score is " ,f1_score(test_Y, predictiona, average = 'weighted'))
```

F1 Score is 0.9079153266165798

b: Data Augmentation: Apply two image augmentation techniques on the Fashion Mnist train data to augment it and then apply the previously developed model on it.

Ans: Model architecture and Hyper-parameters used are same in previous question.

Sample showing Augmented data by rotating images up to 15 degrees, shifting images horizontally and vertically.



Model Summary

```
1 modelb.summary() # Summary of the model |
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_5 (Conv2D)	(None, 28, 28, 32)	320
batch_normalization_5 (Batch Normalization)	(None, 28, 28, 32)	128
max_pooling2d_5 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_6 (Conv2D)	(None, 14, 14, 64)	18496
batch_normalization_6 (Batch Normalization)	(None, 14, 14, 64)	256
max_pooling2d_6 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_7 (Conv2D)	(None, 7, 7, 128)	73856
batch_normalization_7 (Batch Normalization)	(None, 7, 7, 128)	512
max_pooling2d_7 (MaxPooling2D)	(None, 4, 4, 128)	0

conv2d_8 (Conv2D)	(None, 4, 4, 256)	295168
batch_normalization_8 (Batch Normalization)	(None, 4, 4, 256)	1024
max_pooling2d_8 (MaxPooling2D)	(None, 2, 2, 256)	0
conv2d_9 (Conv2D)	(None, 2, 2, 512)	1180160
batch_normalization_9 (Batch Normalization)	(None, 2, 2, 512)	2048
max_pooling2d_9 (MaxPooling2D)	(None, 1, 1, 512)	0
flatten_1 (Flatten)	(None, 512)	0
dense_3 (Dense)	(None, 64)	32832
dropout_2 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 32)	2080
dropout_3 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 10)	330

```

=====
Total params: 1,607,210
Trainable params: 1,605,226
Non-trainable params: 1,984

```

The model has a total of 1,607,210 parameters, out of which 1,605,226 are trainable and 1,984 are non-trainable. The model consists of convolutional layers, batch normalization layers, max pooling layers, fully connected layers with dropout, and an output layer with softmax activation. The architecture progressively reduces the spatial dimensions of the input, from 28x28 to 1x1, while increasing the number of channels in each convolutional layer. The model was compiled with the Adam optimizer (it computes adaptive learning rates for each parameter and performs both momentum and RMSprop style updates for faster convergence), accuracy as the evaluation metric. and categorical cross-entropy loss function, and trained with accuracy as the evaluation metric.

In total, this model has 1,607,210 parameters, out of which 1,605,226 are trainable.

Data augmentation has been done using Keras ImageGenerator. This helps to increase the size of the training dataset by generating new images by applying various transformations like rotation, zooming, shifting. This helps to improve the generalization of the model by reducing overfitting.

```
1 end_timeb = time.time()
```

```
1 total_timeb = end_timeb - start_timeb
```

```
1 print("Total execution time: ", total_timeb, " seconds")
```


Total execution time: 362.29075384140015 seconds


Evaluation

```
1 # evaluating the performance of model on test data
2 test_evalb = modelb.evaluate(test_X, test_Y_one_hot, verbose=0)
```

```
1 print('Test loss:', test_evalb[0])
2 print('Test accuracy:', test_evalb[1])
```

Test loss: 0.34032368659973145
Test accuracy: 0.892799973487854





Confusion Matrix and Performance Score

```
1 cm = confusion_matrix(test_Y, predictionb)
2 #printing the confusion matrix to see correct and wrong predicted values
3 print(cm)
```

```
[[897  0  6 10  1  2 78  0  6  0]
 [ 6 985  0  7  0  0  1  0  1  0]
 [105  1 844  2 25  0 21  0  2  0]
 [122 10  2 823 11  1 23  0  8  0]
 [121  0 74 30 755  0 20  0  0  0]
 [ 0  0  0  0  0  0 994  0  6  0]
 [312  0 52  8 67  0 556  0  5  0]
 [ 5  0  0  0  0 34  0 938  0 23]
 [ 9  1  0  1  0  1  1  0 987  0]
 [ 6  0  0  0  0  8  0 34  0 952]]
```

```
1 print("Precision Score is ", precision_score(test_Y, predictionb, average = 'weighted'))
```

Precision Score is 0.8894451534174008

```
1 print("Recall Score is ", recall_score(test_Y, predictionb, average = 'weighted'))
```

Recall Score is 0.8731

```
1 print("F1 Score is ", f1_score(test_Y, predictionb, average = 'weighted'))
```

F1 Score is 0.874542856446507

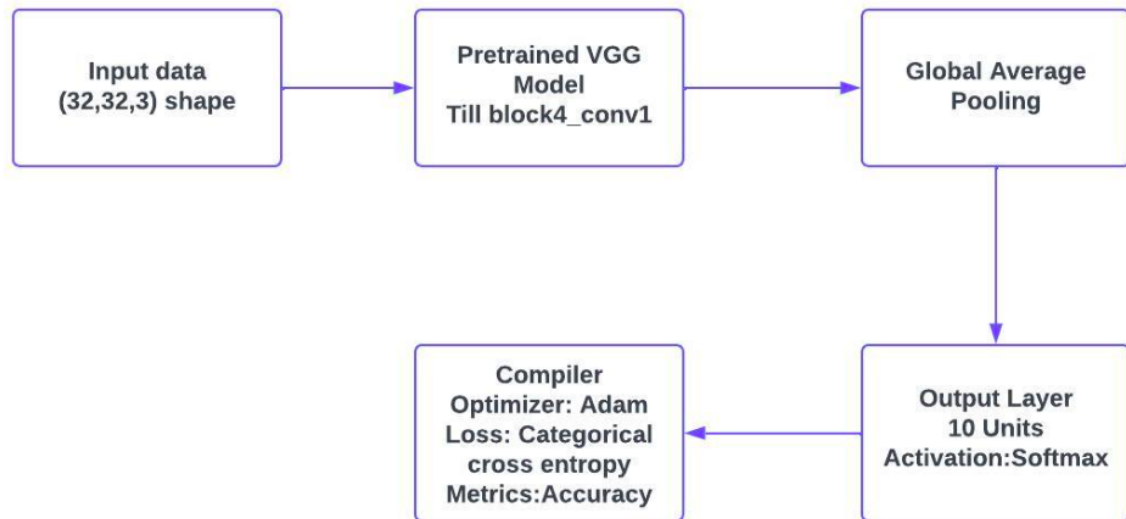
c : Transfer Learning: Load the VGG-19 model. Drop after the block4 conv1 layer (highlighted in the image below) and on top of it add one global average pooling and one final output layer. Keep the base model layers (VGG19) freeze.

Model: "vgg19"

Layer (type)	Output Shape	Param #
input_6 (InputLayer)	[(None, 32, 32, 3)]	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv4 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160
block4_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block4_conv3 (Conv2D)	(None, 4, 4, 512)	2359808

Ans:

Model Architecture



The architecture used on Fashion MNIST dataset starts by reshaping the input image to 28x28x1 and then passing it through a pre-trained VGG network. The model is only used until the conv1 layer of block4. Following this, a Global Average Pooling layer is added to reduce the dimensionality of the feature maps to 512. Finally, a dense output layer is added to predict the class label of the input image.

Hyperparameter in Transfer Learning

- 1) Layers frozen from pretrained model VGG-16= 13
- 2) Layers to tune from pretrained model = 0
- 3) New layers added (Global Average Pooling and Output Layer) = 2

Model Summary

```
1 modelc.summary() # summary of the model
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32, 32, 3)]	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv4 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
dense_8 (Dense)	(None, 10)	5130

Total params: 3,510,858

Trainable params: 5,130

Non-trainable params: 3,505,728

The model is a VGG pretrained model that has been used for classification. It consists of an input layer that takes in an image of shape (32,32,3). The input layer is followed by 2 convolutional layers (block1_conv1 and block1_conv2) with 64 filters each, and a max pooling layer (block1_pool) that reduces the feature map size by half.

Next, there are 2 more convolutional layers (block2_conv1 and block2_conv2) with 128 filters each, followed by another max pooling layer (block2_pool) that reduces the feature map size by half again. The next 4 convolutional layers (block3_conv1, block3_conv2, block3_conv3, and block3_conv4) have 256 filters each, and are followed by a max pooling layer (block3_pool) that reduces the feature map size by half again.

The final convolutional layer (block4_conv1) has 512 filters, and is followed by a global average pooling layer that averages the feature maps to a single vector. This is then followed by a dense layer (dense_8) with 10 units for classification.

The model has a total of 3,510,858 parameters, but only 5,130 of them are trainable, as the convolutional layers are pretrained and frozen.

```
1 end_timec = time.time()
```

```
1 total_timec = end_timec - start_timec
```

```
1 print("Total execution time: ", total_timec, " seconds")
```

Total execution time: 121.30633425712585 seconds

Evaluating the Model

```
1 # evaluating the performance of model in test data
```

```
2
```


```
3 test_evalc = modelc.evaluate(test_images, test_Y_one_hot, verbose=0)
```


```
1 print('Test loss:', test_evalc[0])
```

```
2 print('Test accuracy:', test_evalc[1])
```


Test loss: 0.7428037524223328

Test accuracy: 0.8489000201225281





Confusion Matrix



```
1 print("Precision Score is ", precision_score(test_Y, predictionc, average = 'weighted'))
```

Precision Score is 0.8638325170358898

```
1 print("Recall Score is ", recall_score(test_Y, predictionc, average = 'weighted'))
```

Recall Score is 0.847

```
1 print("F1 Score is ", f1_score(test_Y, predictionc, average = 'weighted'))
```

F1 Score is 0.8445080875402742

d: Make a comparison table including the above three model's performance on the test data (accuracy, number of trainable parameters and execution time).

Ans: All the model are trained for batch size = 32 and epoch = 20

Ques No.	Time Taken in Training	Number of Trainable Parameters	Accuracy	F1 score	Recall	Precision
Ques 1 a	200.49 Seconds	1605226	90.92%	90.79%	90.77%	91.08%
Ques 1 b	362.29 Seconds	1605226	89.27%	87.45%	87.31%	88.94%
Ques 1 c	121.30 Seconds	5130	84.89%	84.45%	84.70%	86.38%

Python 3 Google Compute Engine backend (GPU)

Showing resources from 20:13 to 20:41

System RAM
8.4 / 83.5 GB






GPU RAM
5.4 / 40.0 GB



Disk
25.7 / 166.8 GB







Results

Model comparison was done on the basis of accuracy, training time taken and trainable parameters for all three models i.e. simple convolution model (part a), convolution model with data augmentation (part b) and pretrained vgg model (part c). Trainable parameters were less for part c i.e. 5130 model where other two models have 1605226 trainable parameters. Accuracy is more for part a i.e. simple convolution model (90.92%) whereas other two models have accuracy 89% and 84% respectively. Training time taken was highest for part b where data augmentation was done on the fly i.e. it took around 363 seconds whereas for part c (transfer learning) training time was less i.e. it took around 122 seconds. With reference to accuracy, part a has highest accuracy i.e. 90% whereas accuracy was comparatively less for part b and part c.

Ques 2: Developing ResNet model from scratch

Apply a residual network specified in the following architecture. All convolutional layers use kernel size 3, stride = 1, and padding = "same".


Ans:

Link for Saved Model is given below.

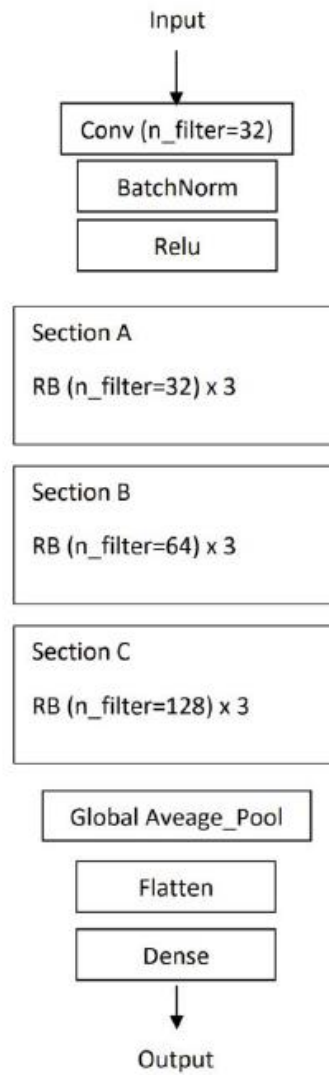
https://drive.google.com/drive/u/2/folders/1anKLUX8RfCPfxEoU0_P16baZzBksvP3g

Model Architecture

Residual Block (RB):



ResNet Structure



The ResNet model architecture consists of the following layers:

Input layer: A 3D tensor of shape (32, 32, 3) representing the input image.

Conv2D layer: A 2D convolutional layer with 32 filters, kernel size of 3, and padding of "same".

Batch Normalization layer: Normalizes the activations of the previous convolutional layer.

Activation layer: Applies the ReLU activation function to the previous layer.

Residual unit: Consists of two 2D convolutional layers with 32 filters, kernel size of 3, and padding of "same", followed by Batch Normalization layers and ReLU activation functions. Concatenates the output of the first layer with the shortcut connection and applies the ReLU activation function.

Residual units: Two more residual units consisting of convolutional, Batch Normalization, and activation layers with 64 and 128 filters, respectively.

Each Residual Block is repeated thrice so that the model can learn more features

Global Average Pooling 2D layer: Computes the spatial average of the previous layer's output across all channels.

Dense layer: A fully connected layer with 10 neurons, representing the output classes.

Softmax activation layer: Applies the softmax function to the output of the previous layer.

Hyperparameter Used

Number of Convolution Layers: 1					
Convolutions Layers	Filter Size	Stride	Number of Filters	Padding	Activation Function
	3x3	1	32	Same	Relu
RSU Block 1					
Convolutions Layers	Filter Size	Stride	Number of Filters	Padding	Activation Function
	3x3	1	32	Same	Relu
Convo Layer 1	3x3	1	32	Same	Relu
RSU Block 2					
Convolutions Layers	Filter Size	Stride	Filters	Padding	Function
	3x3	1	64	Same	Relu
Convo Layer 1	3x3	1	64	Same	Relu
RSU Block 3					
Convolutions Layers	Filter Size	Stride	Filters	Padding	Function
	3x3	1	128	Same	Relu
Convo Layer 1	3x3	1	128	Same	Relu
Output Layer					
Output Layer	Neurons			Activation Function	
	10			Softmax	
Categorical Cross	Adam	0.001	Accuracy		
Training					
Epochs	Batch Size			Validation Split	
10	32			0.2	
Weight Initialiser		Bias Initialiser			
Be default Glorot		By default zeros			

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	(None, 32, 32, 3)	0	[]
conv2d_10 (Conv2D)	(None, 32, 32, 32)	896	['input_2[0][0]']
batch_normalization_10 (Batch Normalization)	(None, 32, 32, 32)	128	['conv2d_10[0][0]']
activation (Activation)	(None, 32, 32, 32)	0	['batch_normalization_10[0][0]']
conv2d_11 (Conv2D)	(None, 32, 32, 32)	9248	['activation[0][0]']
batch_normalization_11 (Batch Normalization)	(None, 32, 32, 32)	128	['conv2d_11[0][0]']
activation_1 (Activation)	(None, 32, 32, 32)	0	['batch_normalization_11[0][0]']
conv2d_12 (Conv2D)	(None, 32, 32, 32)	9248	['activation_1[0][0]']
batch_normalization_12 (Batch Normalization)	(None, 32, 32, 32)	128	['conv2d_12[0][0]']
concatenate (Concatenate)	(None, 32, 32, 64)	0	['activation[0][0]', 'batch_normalization_12[0][0]']
activation_2 (Activation)	(None, 32, 32, 64)	0	['concatenate[0][0]']
conv2d_13 (Conv2D)	(None, 32, 32, 32)	18464	['activation_2[0][0]']
batch_normalization_13 (Batch Normalization)	(None, 32, 32, 32)	128	['conv2d_13[0][0]']
activation_3 (Activation)	(None, 32, 32, 32)	0	['batch_normalization_13[0][0]']
conv2d_15 (Conv2D)	(None, 32, 32, 32)	2080	['activation_2[0][0]']
conv2d_14 (Conv2D)	(None, 32, 32, 32)	9248	['activation_3[0][0]']
batch_normalization_15 (Batch Normalization)	(None, 32, 32, 32)	128	['conv2d_15[0][0]']
batch_normalization_14 (Batch Normalization)	(None, 32, 32, 32)	128	['conv2d_14[0][0]']
concatenate_1 (Concatenate)	(None, 32, 32, 64)	0	['batch_normalization_15[0][0]', 'batch_normalization_14[0][0]']
activation_4 (Activation)	(None, 32, 32, 64)	0	['concatenate_1[0][0]']
conv2d_16 (Conv2D)	(None, 32, 32, 32)	18464	['activation_4[0][0]']
batch_normalization_16 (Batch Normalization)	(None, 32, 32, 32)	128	['conv2d_16[0][0]']
activation_5 (Activation)	(None, 32, 32, 32)	0	['batch_normalization_16[0][0]']
conv2d_18 (Conv2D)	(None, 32, 32, 32)	2080	['activation_4[0][0]']
conv2d_17 (Conv2D)	(None, 32, 32, 32)	9248	['activation_5[0][0]']
batch_normalization_18 (Batch Normalization)	(None, 32, 32, 32)	128	['conv2d_18[0][0]']
batch_normalization_17 (Batch Normalization)	(None, 32, 32, 32)	128	['conv2d_17[0][0]']
concatenate_2 (Concatenate)	(None, 32, 32, 64)	0	['batch_normalization_18[0][0]', 'batch_normalization_17[0][0]']
activation_6 (Activation)	(None, 32, 32, 64)	0	['concatenate_2[0][0]']
conv2d_19 (Conv2D)	(None, 32, 32, 64)	36928	['activation_6[0][0]']

batch_normalization_19 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_19[0][0]']
activation_7 (Activation)	(None, 32, 32, 64)	0	['batch_normalization_19[0][0]']
conv2d_20 (Conv2D)	(None, 32, 32, 64)	36928	['activation_7[0][0]']
batch_normalization_20 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_20[0][0]']
concatenate_3 (Concatenate)	(None, 32, 32, 128)	0	['activation_6[0][0]', 'batch_normalization_20[0][0]']
activation_8 (Activation)	(None, 32, 32, 128)	0	['concatenate_3[0][0]']
conv2d_21 (Conv2D)	(None, 32, 32, 64)	73792	['activation_8[0][0]']
batch_normalization_21 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_21[0][0]']
activation_9 (Activation)	(None, 32, 32, 64)	0	['batch_normalization_21[0][0]']
conv2d_23 (Conv2D)	(None, 32, 32, 64)	8256	['activation_8[0][0]']
conv2d_22 (Conv2D)	(None, 32, 32, 64)	36928	['activation_9[0][0]']
batch_normalization_23 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_23[0][0]']
batch_normalization_22 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_22[0][0]']
concatenate_4 (Concatenate)	(None, 32, 32, 128)	0	['batch_normalization_23[0][0]', 'batch_normalization_22[0][0]']
activation_10 (Activation)	(None, 32, 32, 128)	0	['concatenate_4[0][0]']
conv2d_24 (Conv2D)	(None, 32, 32, 64)	73792	['activation_10[0][0]']
batch_normalization_24 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_24[0][0]']
activation_11 (Activation)	(None, 32, 32, 64)	0	['batch_normalization_24[0][0]']
conv2d_26 (Conv2D)	(None, 32, 32, 64)	8256	['activation_10[0][0]']
conv2d_25 (Conv2D)	(None, 32, 32, 64)	36928	['activation_11[0][0]']
batch_normalization_26 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_26[0][0]']
batch_normalization_25 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_25[0][0]']
concatenate_5 (Concatenate)	(None, 32, 32, 128)	0	['batch_normalization_26[0][0]', 'batch_normalization_25[0][0]']
activation_12 (Activation)	(None, 32, 32, 128)	0	['concatenate_5[0][0]']

conv2d_27 (Conv2D)	(None, 32, 32, 128)	147584	['activation_12[0][0]']
batch_normalization_27 (Batch Normalization)	(None, 32, 32, 128)	512	['conv2d_27[0][0]']
activation_13 (Activation)	(None, 32, 32, 128)	0	['batch_normalization_27[0][0]']
conv2d_28 (Conv2D)	(None, 32, 32, 128)	147584	['activation_13[0][0]']
batch_normalization_28 (Batch Normalization)	(None, 32, 32, 128)	512	['conv2d_28[0][0]']
concatenate_6 (Concatenate)	(None, 32, 32, 256)	0	['activation_12[0][0]', 'batch_normalization_28[0][0]']
activation_14 (Activation)	(None, 32, 32, 256)	0	['concatenate_6[0][0]']
conv2d_29 (Conv2D)	(None, 32, 32, 128)	295040	['activation_14[0][0]']
batch_normalization_29 (Batch Normalization)	(None, 32, 32, 128)	512	['conv2d_29[0][0]']
activation_15 (Activation)	(None, 32, 32, 128)	0	['batch_normalization_29[0][0]']
conv2d_31 (Conv2D)	(None, 32, 32, 128)	32896	['activation_14[0][0]']
conv2d_30 (Conv2D)	(None, 32, 32, 128)	147584	['activation_15[0][0]']
batch_normalization_31 (Batch Normalization)	(None, 32, 32, 128)	512	['conv2d_31[0][0]']
batch_normalization_30 (Batch Normalization)	(None, 32, 32, 128)	512	['conv2d_30[0][0]']
concatenate_7 (Concatenate)	(None, 32, 32, 256)	0	['batch_normalization_31[0][0]', 'batch_normalization_30[0][0]']
activation_16 (Activation)	(None, 32, 32, 256)	0	['concatenate_7[0][0]']
conv2d_32 (Conv2D)	(None, 32, 32, 128)	295040	['activation_16[0][0]']
batch_normalization_32 (Batch Normalization)	(None, 32, 32, 128)	512	['conv2d_32[0][0]']
activation_17 (Activation)	(None, 32, 32, 128)	0	['batch_normalization_32[0][0]']
conv2d_34 (Conv2D)	(None, 32, 32, 128)	32896	['activation_16[0][0]']
conv2d_33 (Conv2D)	(None, 32, 32, 128)	147584	['activation_17[0][0]']
batch_normalization_34 (Batch Normalization)	(None, 32, 32, 128)	512	['conv2d_34[0][0]']
batch_normalization_33 (Batch Normalization)	(None, 32, 32, 128)	512	['conv2d_33[0][0]']
concatenate_8 (Concatenate)	(None, 32, 32, 256)	0	['batch_normalization_34[0][0]', 'batch_normalization_33[0][0]']
activation_18 (Activation)	(None, 32, 32, 256)	0	['concatenate_8[0][0]']
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 256)	0	['activation_18[0][0]']
dense_7 (Dense)	(None, 10)	2570	['global_average_pooling2d_1[0][0]']

=====

Total params: 1,646,858
Trainable params: 1,643,210
Non-trainable params: 3,648

This is a custom-built RESNET model (convolutional neural network model with skip connections). It takes an input of size (32, 32, 3) and has a total of 392,906 parameters, out of which 391,946 are trainable and 960 are non-trainable.

The model has several convolutional layers with batch normalization and activation functions, and skip connections are added between them. It ends with a global average pooling layer followed by a dense layer with 10 units, representing the 10 classes in the dataset.


Overall, the model aims to learn hierarchical representations of the input images through the convolutional layers and use the skip connections to preserve important features from the earlier layers. Finally, the global average pooling layer condenses the learned representations into a fixed-length vector, which is fed to the output dense layer for classification.

Evaluating the Model

```
1 # evaluating the performance of model in test data
2
3 test_evald = model.evaluate(x_testcifar,y_testcifar_one_hot, verbose=0)
```


```
1 print('Test loss:', test_evald[0])
2 print('Test accuracy:', test_evald[1])
```

Test loss: 0.6992190480232239
Test accuracy: 0.763700008392334





Confusion Matrix




```
1 print("Precision Score is ", precision_score(y_testcifar, predictiond, average = 'weighted'))
```

```
Precision Score is  0.7970126970511039
```

```
1 print("Recall Score is ", recall_score(y_testcifar, predictiond, average = 'weighted'))
```

```
Recall Score is  0.7339
```

```
1 print("F1 Score is ", f1_score(y_testcifar, predictiond, average = 'weighted'))
```

```
F1 Score is  0.735701149804657
```

Results

RESNET model has a total of 1,646,858 parameters, out of which 1,643,210 are trainable and 3,648 are non-trainable. Model was trained for 10 epochs and accuracy for the model was 70%. Recall and Precision score for the model is less as compared to above three models. Recall and Precision score is 73% and 79% respectively. This may be due to the complexity of the model as the learning curve shows that training accuracy is more than validation accuracy. Confusion matrix, recall and precision score is given above.